



Keras Cheat Sheet: Neural Networks in Python

April 25th, 2017 in Python



Karlijn Willems

[Keras](#) is an easy-to-use and powerful library for Theano and TensorFlow that provides a high-level neural networks API to develop and evaluate deep learning models.

We recently launched one of the first online interactive deep learning course using Keras 2.0, called "[Deep Learning in Python](#)".

Now, DataCamp has created a Keras cheat sheet for those who have already taken the course and that still want a handy one-page reference or for those who need an extra push to get started.

In no time, this Keras cheat sheet will make you familiar with how you can load datasets from the library itself, preprocess the data, build up a model architecture, and compile, train, and evaluate it. As there is a considerable amount of freedom in how you build up

Furthermore, you'll also see some examples of how to inspect your model, and how you can save ; find examples of how you can predict values for test data and how you can fine tune your models k parameters and early stopping.

learners and start one of our interactive tutorials today!

[Learn R](#)

[Learn Python](#)

In short, you'll see that this cheat sheet not only presents you with the six steps that you can go through to make neural networks in Python with the Keras library.

Python For Data Science Cheat Sheet

Keras

Learn Python for data science interactively at www.datacamp.com

Keras

Keras is a powerful and easy-to-use deep learning library for Theano and TensorFlow that provides a high-level neural networks API to develop and evaluate deep learning models.

A Basic Example

```
>>> import numpy as np
>>> from keras.models import Sequential
>>> from keras.layers import Dense
>>> data = np.random.random((1000,100))
>>> labels = np.random.randint(2,size=(1000,1))
>>> model = Sequential()
>>> model.add(Dense(12,
>>>               activation='relu',
>>>               input_dim=100))
>>> model.add(Dense(1, activation='sigmoid'))
>>> model.compile(optimizer='nnsop',
>>>               loss='binary_crossentropy',
```

Model Architecture

Sequential Model

```
>>> from keras.models import Sequential
>>> model = Sequential()
>>> model.add(Dense(12,
>>>               input_dim=10,
>>>               kernel_initializer='uniform',
>>>               activation='relu'))
>>> model.add(Dense(8, kernel_initializer='uniform', activation='sigmoid'))
```

Multilayer Perceptron (MLP)

Binary Classification

```
>>> from keras.layers import Dense
>>> model.add(Dense(12,
>>>               input_dim=10,
>>>               kernel_initializer='uniform',
>>>               activation='relu'))
>>> model.add(Dense(8, kernel_initializer='uniform', activation='sigmoid'))
```

Multi-Class Classification

```
>>> from keras.layers import Dropout
>>> model.add(Dense(12, activation='relu', input_shape=(784,)))
>>> model.add(Dropout(0.2))
>>> model.add(Dense(10, activation='relu'))
>>> model.add(Dropout(0.2))
>>> model.add(Dense(10, activation='softmax'))
```

Regression

```
>>> model.add(Dense(14, activation='relu', input_dim=train_data.shape[1]))
>>> model.add(Dense(1))
```

Inspect Model

```
>>> model.output_shape
>>> model.summary()
>>> model.get_config()
>>> model.get_weights()
```

Model output shape
Model summary representation
Model configuration
List all weight tensors in the model

Compile Model

MLP Binary Classification

```
>>> model.compile(optimizer='adam',
>>>               loss='binary_crossentropy',
>>>               metrics=['accuracy'])
```

MLP Multi-Class Classification

```
>>> model.compile(optimizer='nnsop',
>>>               loss='categorical_crossentropy',
>>>               metrics=['accuracy'])
```

MLP Regression

```
>>> model.compile(optimizer='nnsop',
>>>               loss='mse',
>>>               metrics=['mse'])
```

Model Training

```
>>> model.fit(x_train,
>>>          y_train,
```

Other

```
>>> from urllib.request import urlopen
>>> data = np.loadtxt(urlopen("http://archive.ics.uci.edu/
>>> ml/machine-learning-database/australian-swine-disease/
>>> data-australian-swine-disease.txt"), delimiter=",")
>>> X = data[:,0:8]
>>> y = data[:,8]
```

Preprocessing

Sequence Padding

```
>>> from keras.preprocessing import sequence
>>> x_train = sequence.pad_sequences(x_train,maxlen=80)
>>> x_test = sequence.pad_sequences(x_test,maxlen=80)
```

One-Hot Encoding

```
>>> from keras.utils import to_categorical
>>> x_train = to_categorical(x_train, num_classes)
>>> x_test = to_categorical(x_test, num_classes)
>>> y_train = to_categorical(y_train, num_classes)
>>> y_test = to_categorical(y_test, num_classes)
```

Train and Test Sets

```
>>> from sklearn.model_selection import train_test_split
>>> x_train,x_test,y_train,y_test = train_test_split(
>>>     x,y,train_size=0.33,
>>>     random_state=0)
```

Standardization/Normalization

```
>>> from sklearn.preprocessing import StandardScaler
>>> scaler = StandardScaler()
>>> x_train = scaler.fit_transform(x_train)
>>> x_test = scaler.transform(x_test)
```

Save / Reload Models

```
>>> from keras.models import load_model
>>> model.save("model.h5")
>>> my_model = load_model("my_model.h5")
```

Model Fine-tuning

Optimization Parameters

```
>>> from keras.optimizers import RMSprop
>>> opt = keras.optimizers.RMSprop(lr=0.001, decay=1e-6)
>>> model2.compile(loss='categorical_crossentropy',
>>>               optimizer=opt,
>>>               metrics=['accuracy'])
```

Early Stopping

```
>>> from keras.callbacks import EarlyStopping
>>> early_stopping_monitor = EarlyStopping(patience=2)
>>> model1.fit(x_train,
>>>          y_train,
>>>          batch_size=32,
>>>          epochs=15,
>>>          validation_data=(x_test,y_test),
>>>          callbacks=[early_stopping_monitor])
```

DataCamp

to any Python for Data Science [interactive](#)

Check Out Keras Cheat Sheet

(Click above to download a printable version or read the online version below.)

learners and start one of our interactive tutorials today!

Python For Data Science Cheat Sheet: Keras

[Learn R](#)[Learn Python](#)

Keras is a powerful and easy-to-use deep learning library for Theano and TensorFlow that provides a high-level neural networks API to develop and evaluate deep learning models.

A Basic Example

```
>>> import numpy as np
>>> from keras.models import Sequential
>>> from keras.layers import Dense
>>> data = np.random.random((1000,100))
>>> labels = np.random.randint(2,size=(1000,1))
>>> model = Sequential()
>>> model.add(Dense(32, activation='relu', input_dim=100))
>>> model.add(Dense(1, activation='sigmoid'))
>>> model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])
>>> model.fit(data,labels,epochs=10,batch_size=32)
>>> predictions = model.predict(data)
```

Data

Your data needs to be stored as NumPy arrays or as a list of NumPy arrays. Ideally, you split the data in training and test sets, for which you can also resort to the `train_test_split` module of `sklearn.cross_validation`.

Keras Data Sets

```
>>> from keras.datasets import boston_housing, mnist, cifar10, imdb
>>> (x_train,y_train),(x_test,y_test) = mnist.load_data()
>>> (x_train2,y_train2),(x_test2,y_test2) = boston_housing.load_data()
>>> (x_train3,y_train3),(x_test3,y_test3) = cifar10.load_data()
```

Other

learners and start one of our interactive tutorials today!

[Learn R](#)

[Learn Python](#)

```
>>> from urllib.request import urlopen
>>> data = np.loadtxt(urlopen("http://archive.ics.uci.edu/ml/machine-learning-databases/pima-indians-diabetes/pima-i
>>> X = data[:,0:8]
>>> y = data[:,8]
```

Preprocessing

Sequence Padding

```
>>> from keras.preprocessing import sequence
>>> x_train4 = sequence.pad_sequences(x_train4,maxlen=80)
>>> x_test4 = sequence.pad_sequences(x_test4,maxlen=80)
```

One-Hot Encoding

```
>>> from keras.utils import to_categorical
>>> Y_train = to_categorical(y_train, num_classes)
>>> Y_test = to_categorical(y_test, num_classes)
>>> Y_train3 = to_categorical(y_train3, num_classes)
>>> Y_test3 = to_categorical(y_test3, num_classes)
```

Train And Test Sets

```
>>> from sklearn.model_selection import train_test_split
>>> X_train5, X_test5, y_train5, y_test5 = train_test_split(X, y, test_size=0.33, random_state=42)
```

Standardization/Normalization

```
>>> standardized_X_test = scaler.transform(x_test2)
```

learners and start one of our interactive tutorials today!

Model Architecture

[Learn R](#)[Learn Python](#)

Sequential Model

```
>>> from keras.models import Sequential
>>> model = Sequential()
>>> model2 = Sequential()
>>> model3 = Sequential()
```

Multi-Layer Perceptron (MLP)

Binary Classification

```
>>> from keras.layers import Dense
>>> model.add(Dense(12, input_dim=8, kernel_initializer='uniform', activation='relu'))
>>> model.add(Dense(8, kernel_initializer='uniform', activation='relu'))
>>> model.add(Dense(1, kernel_initializer='uniform', activation='sigmoid'))
```

Multi-Class Classification

```
>>> from keras.layers import Dropout
>>> model.add(Dense(512, activation='relu', input_shape=(784,)))
>>> model.add(Dropout(0.2))
>>> model.add(Dense(512, activation='relu'))
>>> model.add(Dropout(0.2))
>>> model.add(Dense(10, activation='softmax'))
```

Regression

learners and start one of our interactive tutorials today!

[Learn R](#)

[Learn Python](#)

Convolutional Neural Network (CNN)

```
>>> from keras.layers import Activation, Conv2D, MaxPooling2D, Flatten
>>> model2.add(Conv2D(32, (3,3), padding='same', input_shape=x_train.shape[1:]))
>>> model2.add(Activation('relu'))
>>> model2.add(Conv2D(32, (3,3)))
>>> model2.add(Activation('relu'))
>>> model2.add(MaxPooling2D(pool_size=(2,2)))
>>> model2.add(Dropout(0.25))
>>> model2.add(Conv2D(64, (3,3), padding='same'))
>>> model2.add(Activation('relu'))
>>> model2.add(Conv2D(64, (3, 3)))
>>> model2.add(Activation('relu'))
>>> model2.add(MaxPooling2D(pool_size=(2,2)))
>>> model2.add(Dropout(0.25))
>>> model2.add(Flatten())
>>> model2.add(Dense(512))
>>> model2.add(Activation('relu'))
>>> model2.add(Dropout(0.5))
>>> model2.add(Dense(num_classes))
>>> model2.add(Activation('softmax'))
```

Recurrent Neural Network (RNN)

```
>>> from keras.klayers import Embedding,LSTM
>>> model3.add(Embedding(20000,128))
>>> model3.add(LSTM(128,dropout=0.2,recurrent_dropout=0.2))
>>> model3.add(Dense(1,activation='sigmoid'))
```

Inspect Model

Model output shape

```
>>> model.output_shape
```

learners and start one of our interactive tutorials today!

[Learn R](#)

[Learn Python](#)

Model configuration

```
>>> model.get_config()
```

List all weight tensors in the model

```
>>> model.get_weights()
```

Compile Model

Multi-Layer Perceptron (MLP)

MLP: Binary Classification

```
>>> model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

MLP: Multi-Class Classification

```
>>> model.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['accuracy'])
```

MLP: Regression

```
>>> model.compile(optimizer='rmsprop', loss='mse', metrics=['mae'])
```

learners and start one of our interactive tutorials today!

[Learn R](#)

[Learn Python](#)

Model Training

```
>>> model3.fit(x_train4, y_train4, batch_size=32, epochs=15, verbose=1, validation_data=(x_test4, y_test4))
```

Evaluate Your Model's Performance

```
>>> score = model3.evaluate(x_test, y_test, batch_size=32)
```

Prediction

```
>>> model3.predict(x_test4, batch_size=32)
>>> model3.predict_classes(x_test4, batch_size=32)
```

Save/Reload Models

```
>>> from keras.models import load_model
>>> model3.save('model_file.h5')
>>> my_model = load_model('my_model.h5')
```

Model Fine-Tuning

Optimization Parameters

```
>>> from keras.optimizers import RMSprop
>>> opt = RMSprop(lr=0.0001, decay=1e-6)
>>> model2.compile(loss='categorical_crossentropy', optimizer=opt, metrics=['accuracy'])
```

```
>>> early_stopping_monitor = EarlyStopping(patience=2)
>>> model3.fit(x_train4, y_train4, batch_size=32, epochs=15, validation_data=(x_test4, y_test4), callbacks=[early_st
```

learners and start one of our interactive tutorials today!

[Learn R](#)

[Learn Python](#)

Going Further

Begin with [our Keras tutorial for beginners](#), in which you'll learn in an easy, step-by-step way how to explore and preprocess the wine quality data set, build up a multi-layer perceptron for classification and regression tasks, compile, fit and evaluate the model and fine-tune the model that you have built.

Also, don't miss out on our [Scikit-Learn cheat sheet](#), [NumPy cheat sheet](#) and [Pandas cheat sheet](#)!

What do you think?



Python

Up Next



learners and start one of our interactive tutorials today!

New Course! Supervised Learning in R: Classification

[Learn R](#)

[Learn Python](#)

R Programming

73 views

This beginner-level introduction to machine learning covers four of the most common classification algorithms. You will come away with a b...

September 27th, 2017 in Blog



New Course: Case Studies in Statistical Thinking!

Python

1,135 views

Hone your applied data science skills in Python by doing a variety of case studies across multiple disciplines. Use data science to solve ...

September 20th, 2017 in Blog